

STAT2003 ASSIGNMENT 1

Comparison of Machine Learning Algorithms

Nina Kumagai (19389905) and Amanda Efendi (19488187)

Part 1.1 Design of Experiment

Factors	Types of Machine Learning algorithm
Levels of each factor	5 levels (Logistic Regression, Linear Discriminant Analysis, Classification and Regression Trees, Naïve Bayes, Support Vector Machines)
Experimental Unit	Each prediction run on a test set by a single algorithm trained with the corresponding training set
Response	The accuracy of the predictions made by each algorithm
Number of Replications	10. To provide more than 30 observations across all levels to achieve normality of residuals
Experimental Methodology	Using multiple random sampling seeds, the main dataset is re-sampled into 10 different sets of training and testing sets. Each machine learning algorithm is run on 10 different training sets (10 replicates), and the number of correct predictions is recorded for each replicate. This number is then converted into an accuracy percentage and listed in a longform with the type of machine learning algorithm as their grouping.

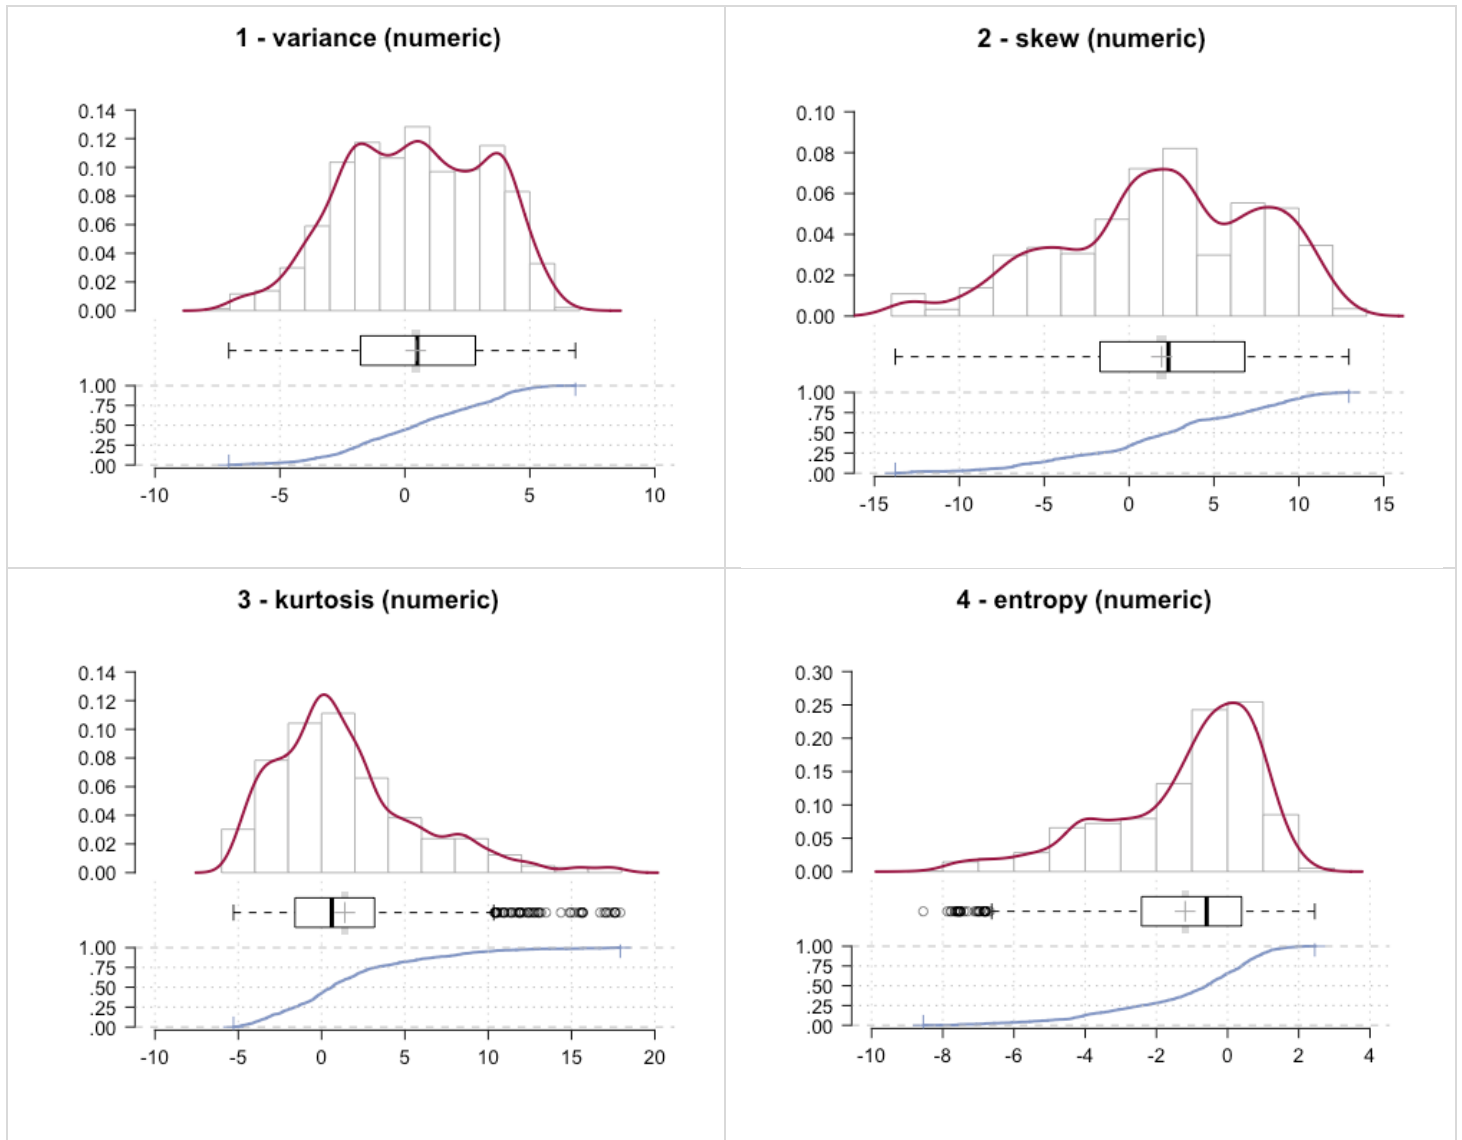
Part 1.2 Principles of Design

Randomisation is applied by the use of the `sample()` function in selecting the training and testing subset. Using the `set.seed()` function, we are able to generate 10 sets of random indexes.

Replication is achieved by generating 10 sets of different training and testing subsets of the data frame. Each Machine Learning algorithms are trained and tested on each set as a replicate. This results in 10 replicates of each machine learning algorithm

Blocking was not applied. According to the UCI site, there were no information on whether or not blocking was used in collecting the banknotes.

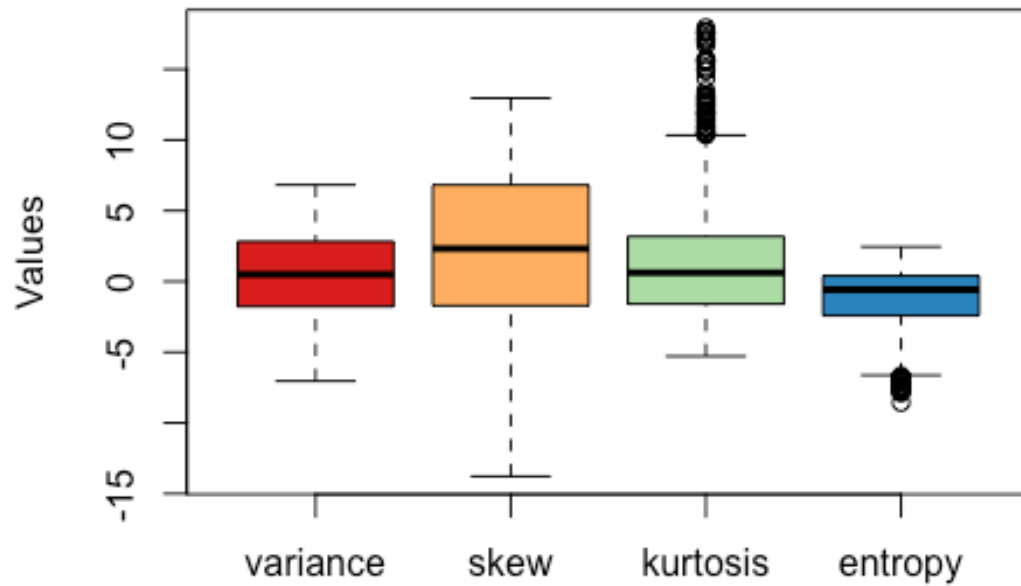
Part 2. EDA Graphs



The banknote dataset has 4 variables along with the classification variable called class. As seen from above, the values for kurtosis and entropy are skewed, while the variance and skew are slightly normal with several peaks. Among the 1372 rows of data, 762 are fake and 610 are real.

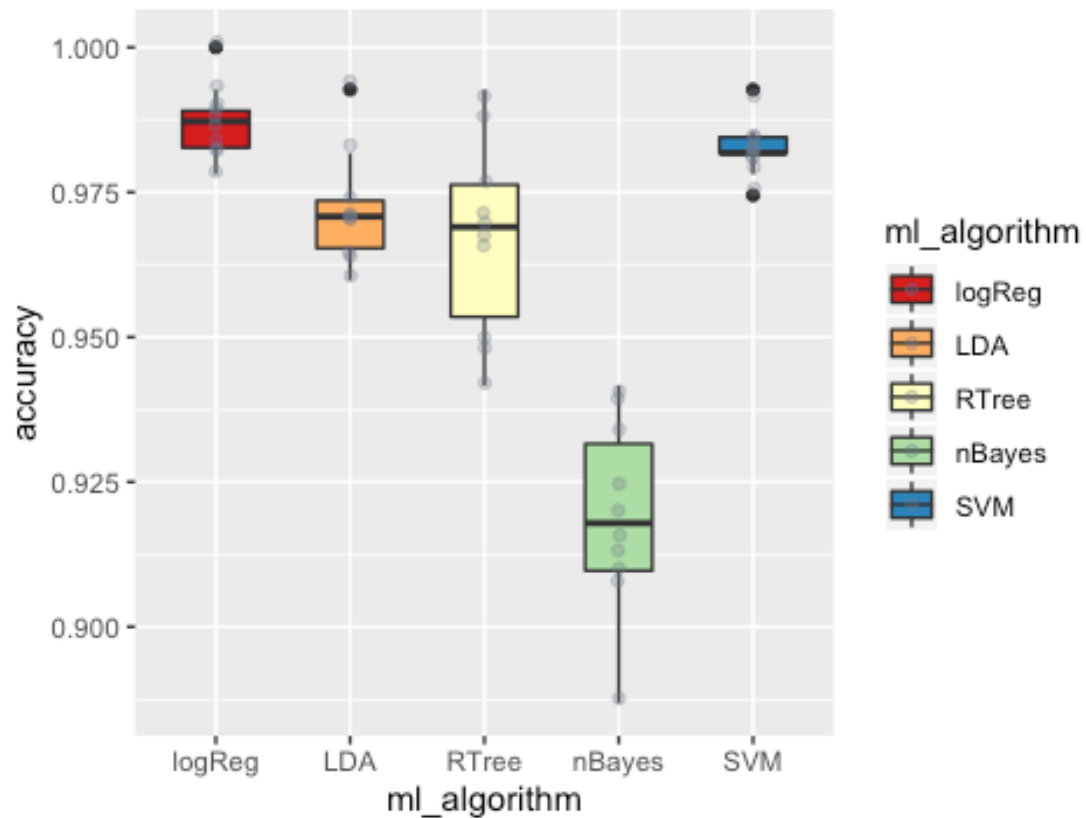
##	variance	skew	kurtosis	entropy
##	Min. :-7.0421	Min. :-13.773	Min. :-5.2861	Min. :-8.5482
##	1st Qu.:-1.7730	1st Qu.: -1.708	1st Qu.: -1.5750	1st Qu.: -2.4135
##	Median : 0.4962	Median : 2.320	Median : 0.6166	Median :-0.5867
##	Mean : 0.4337	Mean : 1.922	Mean : 1.3976	Mean :-1.1917
##	3rd Qu.: 2.8215	3rd Qu.: 6.815	3rd Qu.: 3.1793	3rd Qu.: 0.3948
##	Max. : 6.8248	Max. : 12.952	Max. : 17.9274	Max. : 2.4495

Overview of BankNote Dataset



Part 3. Findings

Boxplot of Machine Learning Algorithm's Accuracy







The boxplot above shows the accuracy of each machine learning algorithm over 10 replicates. The accuracy is obtained by adding up the correct (true negative and true positive) predictions and divided by the total number of observations in the test subset.

LogReg	LDA	RTree	nBayes	SVM
0.9872263	0.9718978	0.9675182	0.9189781	0.9824818



Results of analysis are reported as follows: Mean (standard deviation) of machine learning models for logistic regression, linear discriminant analysis, regression tree, naïve bayes and support vector machine are 0.989(0.006), 0.9777(0.010), 0.9660(0.0115), 0.9124(0.0122), 0.985(0.0085) respectively. As seen on the table above, Logistic Regression and Support Vector Machine have the two highest mean accuracy compared to the other algorithms. The boxplot also indicates that logistic regression and support vector machine algorithm are also consistent with their findings. In contrast, Naïve Bayes, Classification and Regression Trees has a larger variety in their accuracy.


```
##           Df    Sum Sq  Mean Sq F value    Pr(>F)
## ml_algorithm  4 0.03887 0.009717   97.47  <2e-16 ***
## Residuals    45 0.00449 0.000100
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The results of the analysis of variance indicates that at 5% level of significance there is sufficient evidence ($F(4,45)=97.47$, $P= 2 \times 10^{-16}$) to conclude that average accuracy of machine learning models is not the same across model types.



A further examination of the residuals of the ANOVA can be seen below.





Plot of residuals against predicted values does not show any unusual pattern (funneling or bow shape). As normality was valid (Shapiro Wilk, $P=0.7921$), Bartlett test results are reliable. Bartlett test lead to high P-value ($P=0.3147$), indicating equality of variance in residuals across factor levels (homogeneity). Levene's test also leads to the similar result but with slightly higher p-value ($p = 0.5845$). Finally, as the order experiment was random with reference to each factor level setting, independence of errors can be assumed.

##	\$groups
##	accuracy groups
## logReg	0.9872263 a
## SVM	0.9824818 ab
## LDA	0.9718978 bc
## RTree	0.9675182 c
## nBayes	0.9189781 d

For Post-Hoc Comparison Pairwise Test, the Tukey's LSD test shows 4 distinct groups within the levels of treatment. Logistic Regression is considered different to Linear Discriminant Analysis, RTree and Naïve Bayes, while it is the same as Support Vector Machine. While Naïve Bayes, is considered the most distinct from other treatments.

Part 4. Limitations

The bank notes dataset was obtained from the UCI website. There was no information on the process undertaken to collect the data. Thus, randomisation in selecting the bank notes had to be assumed to enable the use of the dataset. The website didn't provide any information regarding any blocking that the owners of the dataset might have done while collecting the data. An example of blocking that could have been done is grouping the notes based on the country of origin.

The banknotes data is then re-sampled into 10 sets of training and testing subsets for the machine learning algorithms. To improve the accuracy of the comparison by assuming normality, more than 10 re-samples of the dataset could have been applied.

Limited understanding of each machine learning algorithms could also impact the interpretation of the results. Perhaps one machine learning algorithm are only more sensitive to large sample sizes, thus resulting in an inaccurate portrayal of its performance in this experiment.

Part 5. Application to other classification problems

As mentioned above, each machine learning algorithm requires different assumption. Logistic Regression performed well in classifying the bank notes, however, if we changed the dataset, it might be a different case. On the other hand, other model may perform better. For example, Naïve Bayes assumes that each variable in a class is independent to all the other variables within that class (Sidana, 2017). In our dataset, the Naïve Bayes performed poorly compared to the other classification algorithms due to the banknote dataset's variables such as variance and skew that are known to be dependent of each other. Due to these reasons, the results from Naïve Bayes will change with a different dataset.

Reference:

Sidana, M. (2017, Feb 28). Types of Classification Algorithms in Machine Learning [Blog post]. Retrieved from <https://medium.com/@Mandysidana/machine-learning-types-of-classification-9497bd4f2e14>.

Appendices

1. Loading the data

```
df = read.table("data_banknote_authentication.txt",
                sep=";",
                col.names=c("variance", "skew", "kurtosis", "entropy", "
class"),
                fill=FALSE,
                strip.white=TRUE)
```

2. Data Wrangling

```
head(df)
```

```
##  variance      skew kurtosis  entropy class
## 1  3.62160    8.6661  -2.8073  -0.44699     0
## 2  4.54590    8.1674  -2.4586  -1.46210     0
## 3  3.86600   -2.6383   1.9242   0.10645     0
## 4  3.45660    9.5228  -4.0112  -3.59440     0
## 5  0.32924   -4.4552   4.5718  -0.98880     0
## 6  4.36840    9.6718  -3.9606  -3.16250     0
```

```
summary(df)
```

```
##      variance              skew              kurtosis              entropy
##  Min.      :-7.0421    Min.      :-13.773    Min.      :-5.2861    Min.      :-8.5482
##  1st Qu.: -1.7730    1st Qu.:  -1.708    1st Qu.: -1.5750    1st Qu.: -2.4135
##  Median :  0.4962    Median :   2.320    Median :  0.6166    Median : -0.5867
##  Mean   :  0.4337    Mean      :   1.922    Mean      :  1.3976    Mean      :-1.1917
##  3rd Qu.:  2.8215    3rd Qu.:   6.815    3rd Qu.:  3.1793    3rd Qu.:  0.3948
##  Max.    :  6.8248    Max.      :  12.952    Max.      :17.9274    Max.      :  2.4495
```

```
##      class
##  Min.      :0.0000
##  1st Qu.: 0.0000
##  Median : 0.0000
##  Mean     : 0.4446
##  3rd Qu.: 1.0000
##  Max.     : 1.0000
```

```
#library(dlookr)
```

```
#describe_data = describe(df)
```

```
#Change class to factor
```

```
df$class <- as.factor(df$class); str(df)
```




```
## 'data.frame': 1372 obs. of 5 variables:
## $ variance: num 3.622 4.546 3.866 3.457 0.329 ...
## $ skew : num 8.67 8.17 -2.64 9.52 -4.46 ...
## $ kurtosis: num -2.81 -2.46 1.92 -4.01 4.57 ...
## $ entropy : num -0.447 -1.462 0.106 -3.594 -0.989 ...
## $ class : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...

#Rename the levels of class
levels(df$class) <- c("Fake", "Real"); str(df)

## 'data.frame': 1372 obs. of 5 variables:
## $ variance: num 3.622 4.546 3.866 3.457 0.329 ...
## $ skew : num 8.67 8.17 -2.64 9.52 -4.46 ...
## $ kurtosis: num -2.81 -2.46 1.92 -4.01 4.57 ...
## $ entropy : num -0.447 -1.462 0.106 -3.594 -0.989 ...
## $ class : Factor w/ 2 levels "Fake", "Real": 1 1 1 1 1 1 1 1 1 1 ...
```


3. EDA on Data

```
library(RColorBrewer)
boxplot(df[,1:4], ylab='Values', main='Overview of BankNote Dataset', col=
c(brewer.pal(4, "Spectral")))
```




```
library(DescTools)
plot(Desc(df))
```


1 - variance (numeric)




2 - skew (numeric)




3 - kurtosis (numeric)



4 - entropy (numeric)

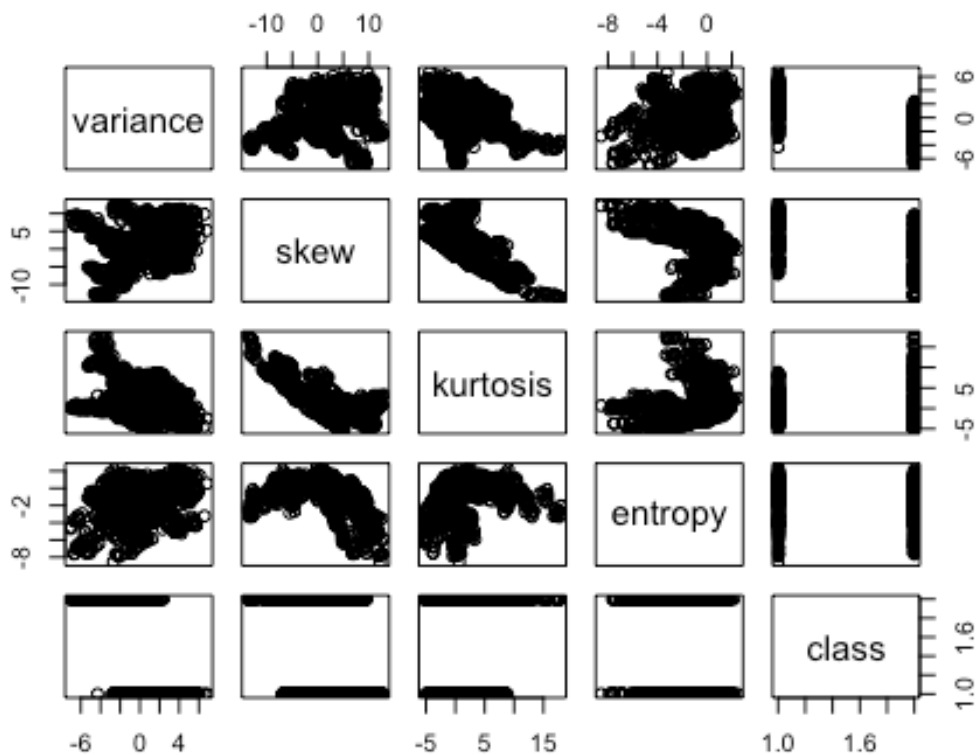


5 - class (factor - dichotomous)



```
#hist(df)
```

```
plot(df)
```



```
library(DataExplorer)
```

4. Making TRAINING and TESTING subsets

```
make_train = function(seed_no){
  set.seed(seed_no)
  testIdx <- sample(1:nrow(df), floor(nrow(df)*0.2))
  training <- df[-testIdx,]
  return (training)
}

make_test = function(seed_no){
  set.seed(seed_no)
  testIdx <- sample(1:nrow(df), floor(nrow(df)*0.2))
  testing <- df[testIdx,]
  return (testing)
}

train1 = make_train(79)
test1 = make_test(79)

train2 = make_train(8)
test2 = make_test(8)

train3 = make_train(19)
test3 = make_test(19)

train4 = make_train(23)
```

```

test4 = make_test(23)

train5 = make_train(32)
test5 = make_test(32)

train6 = make_train(75)
test6 = make_test(75)

train7 = make_train(102)
test7 = make_test(102)

train8 = make_train(421)
test8 = make_test(421)

train9 = make_train(792)
test9 = make_test(792)

train10 = make_train(1)
test10 = make_test(1)

```

5. MACHINE LEARNING ALGORITHMS

TREATMENT: Level 1 Logistic Regression

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```

logreg = function(training, testing){
  glm = glm(class ~ . , data = training, family = binomial(link="logit"))
  # summary(glm)
  pred.glm = predict(glm, newdata=testing, type='response')
  pred.glmclass = rep("Fake", length(pred.glm))
  pred.glmclass[pred.glm>0.5] = "Real"
  # table(pred.glmclass, test1$class, dnn=c("Predictions", "Actual"))
  tn = table(pred.glmclass, testing$class, dnn=c("Predictions", "Actual"))
  [1,1]
  tp = table(pred.glmclass, testing$class, dnn=c("Predictions", "Actual"))
  [2,2]
  accuracy = (tn + tp)/nrow(testing)
  return (accuracy)
}

```

```

logreg_df = data.frame(accuracy = c(logreg(train1, test1), logreg(train2,
test2), logreg(train3, test3), logreg(train4, test4), logreg(train5, test
5), logreg(train6, test6), logreg(train7, test7), logreg(train8, test8), l
ogreg(train9, test9), logreg(train10, test10)),
  ml_algorithm = rep("logReg", 10))

```

```
logreg_df
```

```
##      accuracy ml_algorithm
## 1  0.9854015      logReg
## 2  0.9854015      logReg
## 3  0.9927007      logReg
## 4  1.0000000      logReg
## 5  0.9781022      logReg
## 6  0.9817518      logReg
## 7  0.9890511      logReg
## 8  0.9817518      logReg
## 9  0.9890511      logReg
## 10 0.9890511      logReg
```

```
par(pty='s')
```

```
glm = glm(class ~ . , data = train1, family = binomial(link="logit"))
```


```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
pred.glm = predict(glm, newdata=test1, type='response')
```

```
plot(roc(test1$class, pred.glm), legacy.axes=TRUE)
```

```
## Setting levels: control = Fake, case = Real
```

```
## Setting direction: controls < cases
```



TREATMENT: Level 2 Linear Discriminant Analysis

```
library(MASS)
ldareg <- function(training, testing){
  lda_fit = lda(class ~ . , data=training)
  lda_pred = predict(lda_fit, newdata=testing)
  accuracy = sum(table(testing$class, lda_pred$class)[1,1], table(testing$class, lda_pred$class)[2,2])/nrow(testing)
  return(accuracy)
}


lda_df = data.frame(accuracy = c(ldareg(train1, test1), ldareg(train2, test2), ldareg(train3, test3), ldareg(train4, test4), ldareg(train5, test5), ldareg(train6, test6), ldareg(train7, test7), ldareg(train8, test8), ldareg(train9, test9), ldareg(train10, test10)),
  ml_algorithm = rep("LDA",5))

lda_df
```

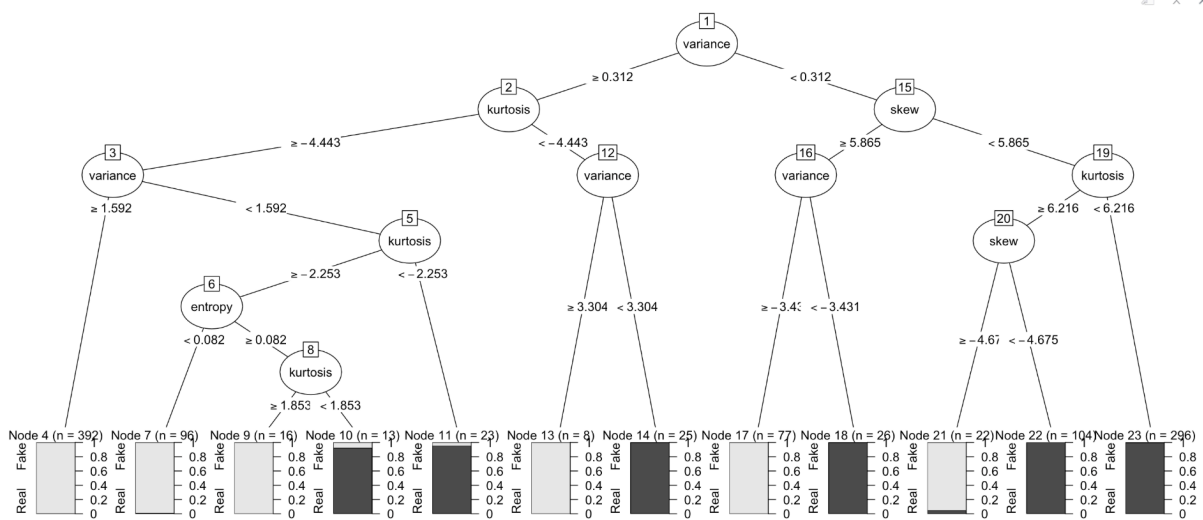
##	accuracy	ml_algorithm
## 1	0.9817518	LDA
## 2	0.9635036	LDA
## 3	0.9927007	LDA
## 4	0.9708029	LDA
## 5	0.9635036	LDA
## 6	0.9598540	LDA
## 7	0.9708029	LDA
## 8	0.9708029	LDA
## 9	0.9708029	LDA
## 10	0.9744526	LDA

TREATMENT: Level 3 Classification and Regression Trees

```
library(rpart)
rpart = rpart(class ~ . , data = train1)
plot(rpart)
text(rpart)
```

```
library(partykit)
plot(as.party(rpart))
```



```
regtree = function(training, testing){
  rpart = rpart(class ~ ., data = training)
  rpart.pred = predict(rpart, newdata = testing, type = "class")
  tn = table(rpart.pred, testing$class, dnn = c("Prediction", "Actual"))
  [1,1]
  tp = table(rpart.pred, testing$class, dnn = c("Prediction", "Actual"))
  [2,2]
  accuracy = (tn + tp)/nrow(testing)
  return (accuracy)
}

regtree_df = data.frame(accuracy = c(regtree(train1, test1), regtree(train2, test2), regtree(train3, test3), regtree(train4, test4), regtree(train5,
```

```
test5), regtree(train6, test6), regtree(train7, test7), regtree(train8, test8), regtree(train9, test9), regtree(train10, test10)),
  ml_algorithm = rep("RTree", 5))
```

regtree_df

```
##      accuracy ml_algorithm
## 1  0.9927007      RTree
## 2  0.9671533      RTree
## 3  0.9416058      RTree
## 4  0.9781022      RTree
## 5  0.9708029      RTree
## 6  0.9489051      RTree
## 7  0.9489051      RTree
## 8  0.9671533      RTree
## 9  0.9708029      RTree
## 10 0.9890511      RTree
```

TREATMENT: Level 4 Naive Bayes

```
library(naivebayes)
```

```
## naivebayes 0.9.6 loaded
```

```
nBayes = function(training, testing){
  nb = naive_bayes(class ~ ., usekernel=T, data=training)
  nb.pred=predict(nb, newdata = testing, type="class")
  tn = table(nb.pred, testing$class, dnn = c("Prediction", "Actual"))[1,1]
  tp = table(nb.pred, testing$class, dnn = c("Prediction", "Actual"))[2,2]
  accuracy = (tn + tp)/nrow(testing)
  return (accuracy)
}
```

```
nbayes_df = data.frame(accuracy = c(nBayes(train1, test1), nBayes(train2, test2), nBayes(train3, test3), nBayes(train4, test4), nBayes(train5, test5), nBayes(train6, test6), nBayes(train7, test7), nBayes(train8, test8), nBayes(train9, test9), nBayes(train10, test10)),
  ml_algorithm = rep("nBayes", 5))
```

nbayes_df

```
##      accuracy ml_algorithm
## 1  0.9197080      nBayes
## 2  0.9233577      nBayes
## 3  0.9160584      nBayes
## 4  0.9416058      nBayes
## 5  0.9379562      nBayes
## 6  0.9343066      nBayes
## 7  0.9124088      nBayes
## 8  0.9087591      nBayes
## 9  0.9087591      nBayes
## 10 0.8868613      nBayes
```

TREATMENT: Level 5 Support Vector Machines

```
# Fitting SVM to the Training set
```

```
library(e1071)
```

```

svm_func = function(train, testing){
  # svm_fit = svm(formula = class ~ ., data = training, type = 'C-classifi
cation', kernel = 'linear')
  training = train
  svm_fit = svm(formula = class ~ ., data = training, kernel = "linear")
  svm.pred = predict(svm_fit, newdata = testing, type = "class")
  tn = table(svm.pred, testing$class, dnn = c("Prediction", "Actual"))[1,
1]
  tp = table(svm.pred, testing$class, dnn = c("Prediction", "Actual"))[2,
2]
  # tn = table(testing[,5], svm_pred)[1,1]
  # tp = table(testing[,5], svm_pred)[2,2]
  accuracy = (tn + tp)/nrow(testing)
  return (accuracy)
}

```

```

svm_df = data.frame(accuracy = c(svm_func(train1, test1), svm_func(train2,
test2), svm_func(train3, test3), svm_func(train4, test4), svm_func(train
5, test5), svm_func(train6, test6), svm_func(train7, test7), svm_func(trai
n8, test8), svm_func(train9, test9), svm_func(train10, test10)),
  ml_algorithm = rep("SVM", 5))

```

svm_df

```

##      accuracy ml_algorithm
## 1  0.9817518      SVM
## 2  0.9817518      SVM
## 3  0.9927007      SVM
## 4  0.9854015      SVM
## 5  0.9744526      SVM
## 6  0.9817518      SVM
## 7  0.9854015      SVM
## 8  0.9817518      SVM
## 9  0.9817518      SVM
## 10 0.9781022      SVM

```

FINAL DATASET OF EACH MODEL'S ACCURACY

```

models_df = rbind(logreg_df, lda_df, regtree_df, nbayes_df, svm_df)
models_df

```

```

##      accuracy ml_algorithm
## 1  0.9854015      logReg
## 2  0.9854015      logReg
## 3  0.9927007      logReg
## 4  1.0000000      logReg
## 5  0.9781022      logReg
## 6  0.9817518      logReg
## 7  0.9890511      logReg
## 8  0.9817518      logReg
## 9  0.9890511      logReg
## 10 0.9890511      logReg
## 11 0.9817518      LDA
## 12 0.9635036      LDA
## 13 0.9927007      LDA

```

```
## 14 0.9708029 LDA
## 15 0.9635036 LDA
## 16 0.9598540 LDA
## 17 0.9708029 LDA
## 18 0.9708029 LDA
## 19 0.9708029 LDA
## 20 0.9744526 LDA
## 21 0.9927007 RTree
## 22 0.9671533 RTree
## 23 0.9416058 RTree
## 24 0.9781022 RTree
## 25 0.9708029 RTree
## 26 0.9489051 RTree
## 27 0.9489051 RTree
## 28 0.9671533 RTree
## 29 0.9708029 RTree
## 30 0.9890511 RTree
## 31 0.9197080 nBayes
## 32 0.9233577 nBayes
## 33 0.9160584 nBayes
## 34 0.9416058 nBayes
## 35 0.9379562 nBayes
## 36 0.9343066 nBayes
## 37 0.9124088 nBayes
## 38 0.9087591 nBayes
## 39 0.9087591 nBayes
## 40 0.8868613 nBayes
## 41 0.9817518 SVM
## 42 0.9817518 SVM
## 43 0.9927007 SVM
## 44 0.9854015 SVM
## 45 0.9744526 SVM
## 46 0.9817518 SVM
## 47 0.9854015 SVM
## 48 0.9817518 SVM
## 49 0.9817518 SVM
## 50 0.9781022 SVM
```

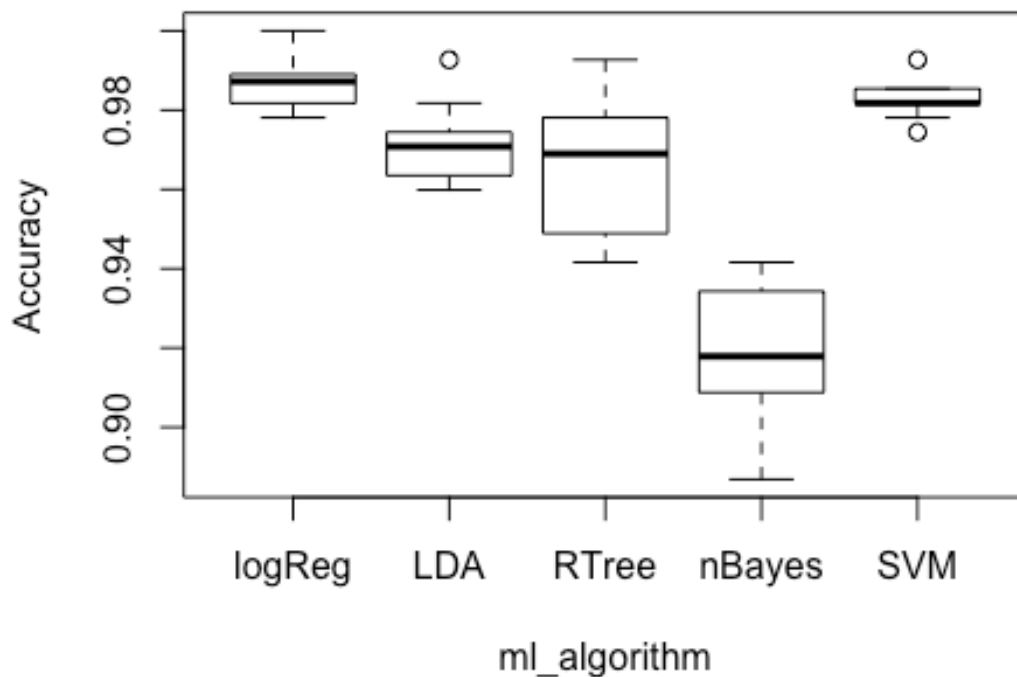
```
# SET DATAFRAME DATATYPE AS NUMERIC AND FACTOR
```

```
models_df$accuracy = as.numeric(models_df$accuracy)
```

```
models_df$ml_algorithm = as.factor(models_df$ml_algorithm)
```

```
boxplot(accuracy~ml_algorithm, xlab="ml_algorithm", ylab="Accuracy", main=  
"Comparison of Accuracy of Machine Learning Models",data=models_df)
```

Comparison of Accuracy of Machine Learning Mod



6. ANOVA Summary

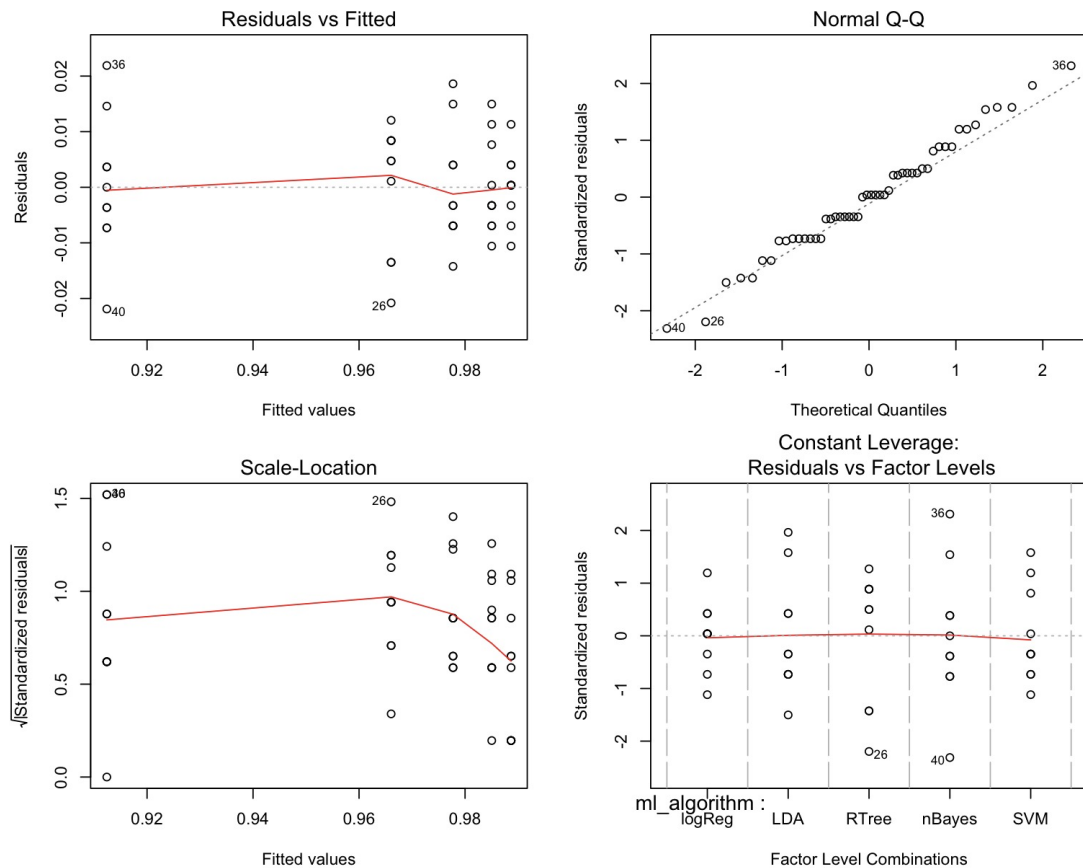
```
ml_anova = aov(accuracy~ml_algorithm,data=models_df)
summary(ml_anova)
```

```
##              Df    Sum Sq  Mean Sq F value    Pr(>F)
## ml_algorithm  4  0.03887  0.009717   97.47  <2e-16 ***
## Residuals    45  0.00449  0.000100
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Conclusion Analysis of variance indicates that at 5% level of significance there is sufficient evidence ($F(4,45)=97.47$, $P= < 2 \times 10^{-16}$) to conclude that average accuracy of machine learning models is not the same across model types.


ANOVA Assumptions

```
opar <- par(mfrow=c(2,2),cex=.8)
plot(ml_anova)
```



```
# follow through from above
# comparing residuals and fitted values of anova
ml_res <- residuals(ml_anova)
ml_pre <- predict(ml_anova)

# Check for Normality of Residuals using Histogram and Boxplot
par(mfrow=c(2,2))
hist(ml_res, xlab="Residuals: ANOVA for Machine Learning Models ", main="Histogram of Residuals")
boxplot(ml_res, ylab="Residuals: ANOVA for Machine Learning Models", main="Boxplot of Residuals")
```




```
shapiro.test(ml_res)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  ml_res
## W = 0.98548, p-value = 0.7921
```

In qqplot not all points are close to the expected line, indicative of some departure from normality and P value for Shapiro Wilks test is high ($P=0.7921$) so there is normality observed here. Next we check for equality of variance of residuals.

```
#Check for equality of variance
plot(ml_res,ml_pre,xlab = "Residuals",ylab = "Predicted Values")
```



```
bartlett.test(accuracy ~ ml_algorithm, data=models_df)

##
## Bartlett test of homogeneity of variances
##
## data: accuracy by ml_algorithm
## Bartlett's K-squared = 4.7431, df = 4, p-value = 0.3147

leveneTest(ml_anova)

## Levene's Test for Homogeneity of Variance (center = median)
##      Df F value Pr(>F)
## group  4 0.7174  0.5845 *
##      45
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Plot of residuals against predicted values does not show any unusual pattern (funneling or bow shape). As normality was valid, Bartlett test results are reliable. Bartlett test lead to high P-value ($P=0.3147$), indicating equality of variance in residuals across factor levels (homogeneity). Levene's test also leads to the similar result but with slightly higher p-value ($p = 0.5845$).

Finally, as the order experiment was random with reference to each factor level setting, independence of errors can be assumed. All underlying anova assumptions are satisfied.

7. Comparing Models

```
tapply(models_df$accuracy,models_df$ml_algorithm, mean)
```

```
##      logReg      LDA      RTree      nBayes      SVM  
## 0.9872263 0.9718978 0.9675182 0.9189781 0.9824818
```

```
tapply(models_df$accuracy,models_df$ml_algorithm, sd)
```

```
##      logReg      LDA      RTree      nBayes      SVM  
## 0.006262549 0.009586807 0.016983770 0.016394044 0.004804968
```

Results of analysis of mean are reported as follows: Mean (standard deviation) of machine learning models for logistic regression, linear discriminant analysis, regression tree, naive bayes and support vector machine are 0.989(0.006), 0.9777(0.010), 0.9660(0.0115), 0.9124(0.0122), 0.985(0.0085) respectively.

8. Multiple Comparisons

Fishers Least Significant Difference Test (Fisher's LSD)

```
library(agricolae)
```

```
MComLSD=LSD.test(ml_anova,"ml_algorithm");MComLSD
```

```
## $statistics  
##      MSerror Df      Mean      CV  t.value      LSD  
## 0.0001422854 45 0.9656204 1.235304 2.014103 0.01074427  
##  
## $parameters  
##      test p.adjusted      name.t ntr alpha  
## Fisher-LSD      none ml_algorithm 5 0.05  
##  
## $means  
##      accuracy      std  r      LCL      UCL      Min      Max  
## LDA 0.9718978 0.009586807 10 0.9643005 0.9794952 0.9598540 0.9927007  
## logReg 0.9872263 0.006262549 10 0.9796289 0.9948236 0.9781022 1.0000000  
## nBayes 0.9189781 0.016394044 10 0.9113808 0.9265754 0.8868613 0.9416058  
## RTree 0.9675182 0.016983770 10 0.9599209 0.9751156 0.9416058 0.9927007  
## SVM 0.9824818 0.004804968 10 0.9748844 0.9900791 0.9744526 0.9927007  
##      Q25      Q50      Q75  
## LDA 0.9653285 0.9708029 0.9735401  
## logReg 0.9826642 0.9872263 0.9890511  
## nBayes 0.9096715 0.9178832 0.9315693  
## RTree 0.9534672 0.9689781 0.9762774  
## SVM 0.9817518 0.9817518 0.9844891  
##  
## $comparison  
## NULL  
##  
## $groups  
##      accuracy groups  
## logReg 0.9872263 a  
## SVM 0.9824818 ab  
## LDA 0.9718978 bc  
## RTree 0.9675182 c  
## nBayes 0.9189781 d
```

```
##
## attr(,"class")
## [1] "group"
```

Tukey's Studentised Range Test

```
MComTukey=HSD.test(ml_anova,"ml_algorithm");MComTukey
```

```
## $statistics
##      MSerror Df      Mean      CV      MSD
## 0.0001422854 45 0.9656204 1.235304 0.01515777
##
## $parameters
##      test      name.t ntr StudentizedRange alpha
## Tukey ml_algorithm    5          4.018417 0.05
##
## $means
##      accuracy      std  r      Min      Max      Q25      Q50
## LDA      0.9718978 0.009586807 10 0.9598540 0.9927007 0.9653285 0.9708029
## logReg 0.9872263 0.006262549 10 0.9781022 1.0000000 0.9826642 0.9872263
## nBayes 0.9189781 0.016394044 10 0.8868613 0.9416058 0.9096715 0.9178832
## RTree 0.9675182 0.016983770 10 0.9416058 0.9927007 0.9534672 0.9689781
## SVM      0.9824818 0.004804968 10 0.9744526 0.9927007 0.9817518 0.9817518
##
##      Q75
## LDA      0.9735401
## logReg 0.9890511
## nBayes 0.9315693
## RTree 0.9762774
## SVM      0.9844891
##
## $comparison
## NULL
##
## $groups
##      accuracy groups
## logReg 0.9872263      a
## SVM      0.9824818      ab
## LDA      0.9718978      b
## RTree 0.9675182      b
## nBayes 0.9189781      c
##
## attr(,"class")
## [1] "group"
```

Duncan's Test

```
#Use first treatment(alphanumerically)as control
```

```
library(multcomp)
```

```
MComScheffe=glht(ml_anova,Linfct=mcp(Treatment="Dunnett"))
summary(MComScheffe)
```

```
##
## Simultaneous Tests for General Linear Hypotheses
##
## Fit: aov(formula = accuracy ~ ml_algorithm, data = models_df)
```

```
##
## Linear Hypotheses:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) == 0      0.987226   0.003772 261.720 < 0.001 ***
## ml_algorithmLDA == 0  -0.015328   0.005335  -2.873  0.02315 *
## ml_algorithmRTree == 0 -0.019708   0.005335  -3.694  0.00252 **
## ml_algorithmnBayes == 0 -0.068248   0.005335 -12.794 < 0.001 ***
## ml_algorithmSVM == 0  -0.004745   0.005335  -0.889  0.80710
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## (Adjusted p values reported -- single-step method)
```

Pairwise T-tests using Bonferroni and Holm

```
accuracy=models_df$accuracy
ml_algorithm=models_df$ml_algorithm

MComBonferroni=pairwise.t.test(accuracy,ml_algorithm,p.adjust="bonferroni")
MComBonferroni

##
## Pairwise comparisons using t tests with pooled SD
##
## data:  accuracy and ml_algorithm
##
##      logReg  LDA      RTree   nBayes
## LDA      0.0618 -        -        -
## RTree     0.0059 1.0000 -        -
## nBayes    1.4e-15 6.7e-12 9.2e-11 -
## SVM       1.0000 0.5337 0.0740 1.7e-14
##
## P value adjustment method: bonferroni

attach(models_df)
MComPairwise=pairwise.t.test(accuracy,ml_algorithm);MComPairwise

##
## Pairwise comparisons using t tests with pooled SD
##
## data:  accuracy and ml_algorithm
##
##      logReg  LDA      RTree   nBayes
## LDA      0.0309 -        -        -
## RTree     0.0036 0.7570 -        -
## nBayes    1.4e-15 5.3e-12 6.4e-11 -
## SVM       0.7570 0.1601 0.0309 1.5e-14
##
## P value adjustment method: holm
```

RESULT PLOTTING

```
library(ggplot2)
p = ggplot(models_df, aes(ml_algorithm, accuracy, fill=ml_algorithm))
p + geom_boxplot(width=0.5) + geom_jitter(width = 0.01, colour = 'lightsteelblue4', alpha=0.3) + scale_fill_brewer(palette = "Spectral") + ggtitle("Boxplot of Machine Learning Algorithm's Accuracy")
```

Boxplot of Machine Learning Algorithm's Accuracy

